

03.3 Structural Patterns

Eight Patterns and Their Framework Design Implications

Eight structural patterns were identified in the Stage 1 evidence base. Each pattern is presented with a key claim, the supporting evidence (drawing on quantitative, qualitative, and policy records as available), and the framework design decision it motivated. The patterns are partitioned into two groups by the analytic basis on which they were detected. Group A patterns are observable across the full evidence base independent of cross-domain comparison; they describe asymmetries within or between tool codes that single-tool analyses can surface. Group B patterns emerge only through systematic cross-domain mapping enabled by the standardised extraction structure; they describe relationships across populations, delivery parameters, and finding directions that are visible only when records can be aggregated across categories. The main manuscript summarises these eight patterns as four: translation gaps (Patterns 1 to 3), population misalignment (Pattern 4), reporting limitations (Pattern 5), and design contraindications (Patterns 6 and 8); Pattern 7 (self-regulatory strategies as design resources) is reported here only.

All counts derive from the taxonomy dataset (n = 410 records). Tool IDs (T001–T210) reference entries in the RG Intervention Framework. Direction codes (positive, mixed, negative, null) and quality tiers (HIGH, MODERATE-HIGH, MODERATE, LOW-MODERATE, LOW) are defined in 02.4 Evidence Direction and Sample Size and 02.5 Evidence Quality and Readiness. Tool codes are defined in 02.3 Tool Codes and Domains.

Group A. Patterns observable across the full evidence base

Pattern 1. Laboratory-to-real-world translation: divergence (MSG-POP) and convergence (BREAK-FORCE)

Same phenomenon, opposite outcomes: laboratory evidence systematically overstates MSG-POP effectiveness, while BREAK-FORCE laboratory predictions converge with deployed dose-response data and supply adaptive calibration logic.

MSG-POP shows divergence. Of 33 MSG-POP records, 24 were conducted in laboratory settings with simulated electronic gaming machines, with median sample n = 107 and direction ratings predominantly positive. The 9 records from live gambling environments show a different pattern: at one large operator, 0.67% of high-intensity sessions terminated at the pop-up threshold under standard messaging, rising to 1.39% under enhanced messaging (Auer et al., 2015). Online experimental records showed no immediate effects on wagered amounts or session duration; significant reductions appeared only at 15-day follow-up, a pattern inconsistent with the behavioural-suppression mechanism tested in laboratory designs (Caillon et al., 2021). Qualitative records explain the mechanism: warning language framing recipients as problem gamblers activates defensiveness and label rejection, while content emphasising spending awareness is perceived as informative and non-stigmatising (Hing et al., 2019; Gainsbury et al., 2018). Static-message regulatory mandates have preceded behavioural evidence of effectiveness in multiple jurisdictions: UK signage compliance increased from 4.1% pre-Gambling-Act to 65% post-Act regardless of evidence of behavioural impact (Moodie & Reith, 2009), and a French operator audit found 100% compliance with mandated banner content despite messages being noted as easily missed at the top of the page (Marionneau et al., 2017). Institutional rationales for deployment are thus established independently of demonstrated impact.

BREAK-FORCE shows convergence. Laboratory work established that active breaks requiring a behavioural response are more effective than passive pauses (Blaszczynski et al., 2016; Parke et al., 2022). Norwegian deployment showed 90-second breaks insufficient to achieve responsible-

gambling goals (Auer et al., 2019). Austrian deployment showed 5-minute breaks producing modest effects (median time-to-next-session 1.6 vs 0.6 min control) and 15-minute breaks producing substantially larger reductions (Hopfgartner et al., 2022). A deposit-frequency-triggered 60-minute break produced same-day reductions, adding a trigger-condition dimension to the dose question (Auer & Griffiths, 2023). Schwarzkopf et al.'s (2024) laboratory finding that arousal reduction mediates break effectiveness supplies the mechanism connecting these dose findings. Evidence spanning eight years, four countries, three platform types, and two study designs converges on a coherent dose-response model: break effectiveness increases with duration, requires sufficient length to interrupt autonomic arousal, and is enhanced by active rather than passive engagement during the break.

Framework response. The Framework calibrates BREAK-FORCE specifications from the Norsk Tipping deployed dose-response study (T030 90s marked minimal-effect; T031 5min and T032 15min carrying deployed effect data). Trigger condition is treated as a separate design parameter, with deposit-frequency triggering (T037, ComeOn) catalogued distinctly from duration-triggered breaks. MSG-POP entries (T058–T071) are dominated by laboratory-tested designs but are accompanied by extrapolated entries (T191, T192, T195, T196, T208) that translate qualitative and mechanistic evidence into platform-deployable variants the deployed research has not yet covered.

Primary domains. MSG-POP, BREAK-FORCE, STRUCT.

Pattern 2. Untranslated evidence clusters: undeployed detection knowledge and clinical content

Two inverse translation failures: precise actionable knowledge without intervention application (DETECT, GUIDE-LRGL) and validated intervention content without a deployment pathway (THER cluster).

Undeployed knowledge: DETECT and GUIDE-LRGL. DETECT (n = 11; 9 Conceptual, 2 Deployed) contains the most quantitatively precise harm-marker evidence in the taxonomy. Within-session deposit frequency emerges as one of the strongest single behavioural markers: high-risk gamblers deposit at 3.6× the rate of low-risk gamblers (Auer et al., 2023); first-week behavioural data alone predicts 90-day high-risk status (Auer et al., 2024); slot machine play shows the strongest associations with behavioural harm markers across gambling verticals (Delfabbro et al., 2023). The two Deployed records describe operator-side detection systems but do not connect a validated marker to a triggered intervention. GUIDE-LRGL (n = 13; all Conceptual) carries thresholds derived across multiple Canadian national samples and international replications: ≤8 occasions/month and ≤\$75 CAD/month (approximately 1–2% of household income) predict harm onset over five years (Currie et al., 2021; Young et al., 2022). Activity-specific thresholds with acceptable discrimination (AUC ≥ 0.70) exist for eight gambling verticals (Dowling et al., 2022); income moderation means low-income gamblers reach harm-associated consumption at substantially lower absolute spending (Langeland et al., 2022). GUIDE-LRGL has the highest concentration of MODERATE-HIGH evidence in the taxonomy (3 of 13, 23.1%). No record in the taxonomy uses a validated behavioural marker to trigger an intervention, or a validated consumption threshold to calibrate a limit default.

Undeployed clinical content: the THER cluster. The 58 records with primary THER-* tools constitute the methodologically strongest cluster in the taxonomy by quality distribution. THER-CBT (n = 15) yielded 12 Validated records with positive effects on PGSI, gambling frequency, and financial outcomes; THER-MI (n = 12) produced 10 EXP records with consistent positive effects; THER-COMBO (n = 10), all EXP, includes the 2008 internet-delivered CBT trial that established online delivery feasibility (Carlbring & Smit, 2008). Across all 58 records, only 5 are Deployed. The deployment gap reflects delivery context, not evidence quality: clinical and research-setting formats (telephone sessions, structured face-to-face programmes, online workbooks requiring sustained engagement) do not map onto the frictionless, automated delivery architecture of digital gambling

platforms. Qualitative records add a substantive constraint: a transcript analysis of 85 web-based counselling sessions found gamblers split between crisis presentation (49%) requiring emotional containment and strategy-seeking presentation (51%) requiring structured skill-building, with both served by the same MI-derived protocol (Rodda et al., 2015). CBT exposure components were transformative for abstinence-oriented gamblers but produced early dropout among those pursuing controlled gambling, an incompatibility between exposure paradigms and harm-reduction goals (Smith et al., 2016). Validated therapeutic mechanisms are technique-specific and presentation-specific in ways a single automated tool cannot accommodate.

Framework response. DETECT entries (T090–T096) carry the rule-based, ML, and trajectory-prediction algorithms identified in the taxonomy. Because none of these algorithms had been connected to a player-facing tool in the deployed evidence, they enter the Framework as detection components rather than standalone tools, and two extrapolated entries (T209 response-latency dissociation; T210 emotional valence trajectory) operationalise harm-marker thinking that the deployed evidence has not extended to. GUIDE-LRGL guidelines (T162–T166) supply default values for limit calibration; T167 (operator-implemented LRGL defaults, theoretically derived) directly addresses the translational gap the taxonomy made visible. THER content (T121–T136) is retained for reference; the translatable subset is THER-SELF and brief MI (T123, T126, T127, T129, T132), which underpin self-directed digital prevention without clinical infrastructure. The qualitative crisis-vs-strategy distinction informs the tool-selection logic: directive MI techniques are suitable for moderate-risk non-treatment-seeking users only, with crisis-presenting users routed to CONTACT-PROACTIVE entries.

Primary domains. DETECT, GUIDE-LRGL, THER-SELF, THER-MI, THER-COMBO.

Pattern 3. Platform-integrated but platform-untested: structural modifications

The structural modifications most strongly associated with online harm (auto-play, frictionless payment, continuous session) are documented as escalation drivers rather than as targets of evaluated protective modification; STRUCT-STAKE is uniquely the tool type where the highest-risk users show the most favourable response.

Structural modifications operate at the product-design level. When evaluated collectively across 33 records (primary STRUCT-* tool), risk-level response patterns emerge that are not visible in domain-level summaries. STRUCT-STAKE consistently shows the strongest expenditure effects among moderate- and high-risk gamblers, making it the only tool type in the taxonomy where the highest-risk users exhibit the most favourable response to modification (Błaszczynski et al., 2005; Sharpe et al., 2005; Jackson et al., 2016). STRUCT-DISPLAY records show effects concentrated in at-risk rather than high-risk gamblers. STRUCT-OTHER records diverge by player demand: modifications that actively interrupt the betting sequence yield inhibitory effects with acceptable enjoyment ratings (Harris et al., 2018), while those introducing player-controlled automation or skill-attribution features result in escalation. Auto-play introduction at a live operator increased wagered amounts by 6.9–9.1% (Jonsson et al., 2024), the clearest example in the taxonomy of a deployed structural feature with a measurable population-level harm-increasing effect.

The platform-untested gap. Of 33 STRUCT records, 29 are from laboratory or land-based venues. Only 4 records have an online platform setting: 1 Delphi consensus, 1 qualitative focus group, 1 content audit, and 1 single-session laboratory experiment in online context. No record evaluates a protective structural modification in a live online environment using objective behavioural outcomes. The structural features most strongly associated with online harm (auto-play, cashless frictionless payment, continuous session architecture) are documented in the taxonomy almost exclusively as causes of escalation, not as targets of evaluated protective modification. A qualitative online record

found participants describing payment design as a tool for circumventing self-exclusion rather than supporting it (Lakew et al., 2024).

Framework response. The Framework retains structural modifications across multiple sub-categories but treats most as carrying evidence drawn from non-online contexts. STRUCT-TEMPO entries (T151–T153) carry the inter-spin delay evidence, with T205 extrapolating the BOI 1700ms threshold finding into a session-risk-responsive adaptive variant rather than a universal lock. STRUCT-PAY combines policy-level entries (T155–T158) with extrapolated digital tools (T178 reversal friction, T180 one-click withdrawal, T199 cross-operator budget cap) that translate land-based payment-friction logic into the online context. STRUCT-OTHER (T150 auto-play removal, T160 demo-RTP parity) carries the policy-derived modifications, and STRUCT-DISPLAY entries (T147–T149, T206) supply the persistent salience-cue tools that have laboratory support but minimal deployed online evaluation.

Primary domains. STRUCT-TEMPO, STRUCT-STAKE, STRUCT-PAY, STRUCT-OTHER, STRUCT-DISPLAY.

Group B. Patterns apparent only through cross-domain mapping

Pattern 4. Population misalignment: the prevention paradox in evidence and uptake

Two compounding mismatches: research focuses on high-risk populations while voluntary tool uptake reaches only those already aware of their risk, leaving moderate-risk gamblers underserved by both evidence and tools.

The prevention paradox in evidence focus. Tool target tier classifications across 410 records show 157 (38.3%) targeting HIGH-risk gamblers, 90 (22.0%) MODERATE, 79 (19.3%) LOW, 54 (13.2%) ALL, and 30 (7.3%) not applicable to player-facing targeting. The LOW-risk count is inflated by SELF-REG (n = 30, all Conceptual) and GUIDE-LRGL (n = 13, all Conceptual), neither of which corresponds to a platform-embedded experimentally-evaluated tool. Excluding these 43 Conceptual records, the experimentally-evaluated platform-integrated evidence base for low-risk gamblers is minimal. Among the 46 EXP extractions targeting low-risk populations, 37 were laboratory studies evaluating messaging or structural EGM modifications using simulated tasks. This pattern reflects a clinical orientation: research focuses on high-risk individuals while moderate- and low-risk gamblers, who contribute most to aggregate gambling harm, remain underserved by both deployed tools and experimental evidence.

The voluntary uptake paradox. Observational LIMIT-DEP records consistently report voluntary deposit limit uptake at 0.3–1.31% of active accounts (Broda et al., 2008; Auer et al., 2020), with between-operator range 1.2–8.5% in a UKGC audit (Heirene et al., 2021), increasing to 8.3% when integrated with a loyalty programme (Auer et al., 2021). Limit-setters show higher baseline gambling intensity than non-setters; problem gamblers are significantly more likely to use limits than low-risk gamblers (Heirene et al., 2021). Voluntary limits are adopted primarily by users already aware of their elevated risk, leaving moderate-risk gamblers, those who would benefit most from default or prompted limits, unaddressed. Qualitative records explain why: gamblers distinguish self-initiated limits (perceived as personal control) from operator-prompted limits (perceived as surveillance or pathology attribution), reducing adoption among moderate-risk users who do not identify with the problem-gambling label (Wood & Griffiths, 2008; Hing et al., 2019). Prompt timing matters: prompts at account creation or login (cool-state decision points) are associated with higher uptake than during-session prompts. Policy records document mandatory deposit-limit defaults under the Netherlands KOA Act (2021), establishing regulatory infrastructure for default-setting logic.

Framework response. The Framework operationalises tier targeting through the Target Population field on each tool rather than through duplicate per-tier entries. Universal-target tools dominate the Framework, reflecting the framework principle that most framework tools should be available across all tiers with delivery defaults shifting by risk. Tier-specific entries cluster at the high-risk end (THER cluster T123–T136; CONTACT-PROACTIVE T097–T101) where the literature concentrates. The moderate-risk uptake gap shaped a distinct group of default-on and risk-triggered entries: T010 opt-out registration prompt (Empirically tested), T011 algorithmically-suggested deposit limit (Theoretically derived), and the extrapolated T202 risk-adaptive dynamic default limits and T204 risk-score-linked ceilings. These translate the prevention-paradox finding into default-on or risk-triggered architectures rather than relying on voluntary uptake. The autonomy-supportive vs action-oriented framing distinction sits in tool delivery specifications rather than as separate entries.

Primary domains. LIMIT-DEP, SELF-REG, GUIDE-LRGL, CONTACT-PROACTIVE, MSG-POP.

Pattern 5. Systematic reporting gaps: delivery timing and mechanism specification

Missing data treated as evidence: blank extraction fields for delivery timing, behaviour change technique specification, and personalisation delineate the limits of what the published literature currently supports.

Delivery timing. Component-level analysis of delivery timing across 410 records shows DURING-PLAY 113 (27.6%), ON-DEMAND 100 (24.4%), PRE-SESSION 83 (20.2%), CONTINUOUS 58 (14.1%), BETWEEN-SESSIONS 40 (9.8%), SCHEDULED 38 (9.3%), POST-SESSION 14 (3.4%), and not applicable 11 (2.7%). Each delivery phase maps to distinct psychological states and implementation requirements. Pre-play delivery is most common for limit-setting and self-assessment tools, where session arousal is absent and uptake is facilitated. During-play delivery dominates for structural modifications, messaging, and feedback. Post-session delivery is markedly underrepresented (3.4%) despite theoretical rationale: lower arousal, absence of active financial stakes, restored deliberative capacity. Account-level delivery (limit defaults, long-term self-exclusion, risk-based proactive outreach) is captured by the BETWEEN-SESSIONS and SCHEDULED categories.

Mechanism and content specification. Tool_Content and Delivery_Mechanism extraction fields show systematic shortcomings in the primary literature. Most MSG-POP records that include verbatim content describe message text from laboratory simulations; real-world operator pop-up content is documented in only a minority of Deployed records. Delivery mechanism specification varies considerably: some records specify trigger condition, frequency, duration, dismissal logic, and content comprehensively; many specify only modality and timing phase. Behaviour change technique specification is absent from most records and, where present, is typically inferred from theoretical framing rather than pre-specified. Personalisation distribution shows NONE 166 (40.5%), BEHAVIOURAL 109 (26.6%), USER-SET 102 (24.9%), INDIVIDUALISED 36 (8.8%), not applicable 6 (1.5%). The combined behavioural and individualised personalisation share (35.4%) reflects active research on adaptive content; identical-content delivery still characterises 40% of records.

Missing data as a data point. The taxonomy preserves extraction fields that were frequently impossible to complete. Blank cells in delivery mechanism, tool content, and BCT-specification columns are not omissions to be filled but data points in their own right, delineating the limits of what the published literature supports for cross-study synthesis of intervention logic. The structure of the missingness identifies where the field needs to standardise reporting before further synthesis is feasible.

Framework response. The Framework makes timing and personalisation explicit specification fields on every tool: each entry declares its delivery setting, trigger mechanism, and personalisation type, surfacing parameters that the taxonomy showed are inconsistently reported. Post-session delivery, underrepresented in the evidence base, is operationalised through two extrapolated entries (T193 post-session wellbeing debrief, T210 cross-session emotional valence trajectory) that occupy a delivery window the taxonomy left underspecified. On personalisation, behavioural and adaptive variants account for 82 of 210 Framework tools (39%), with adaptive entries including T011, T064, T175, T176, T182, T202–T204 and the personalised inter-spin delay T205. These extrapolated tools translate the adaptive-personalisation ambition into platform-deployable specifications drawing on detection signals, session state, and trajectory-based prediction.

Primary domains. MSG-POP, FEED-PERSONAL, LIMIT-DEP, BREAK-FORCE, CONTACT-PROACTIVE.

Pattern 6. The for-whom question: moderation evidence and risk-level reversals

The same intervention can be protective, null, or harm-increasing across risk tiers, gender, and other moderators; for-whom is a safety question, not just an optimisation question, and aggregate effect summaries hide the reversals.

Risk-level reversals. Where subgroup analyses are reported, the assumption that interventions are most effective for highest-risk players is not consistently supported and in several cases is reversed. EDU-VIDEO active-learning content increased limit-setting intentions for low-risk ($b = 1.21$, $p = .004$) and moderate-risk ($b = 0.76$, $p = .03$) gamblers but had no effect on high-risk gamblers (Hollingshead et al., 2019). Newall et al. (2023) reported a credible PGSI $\times$ message interaction in which the message+ condition amplified the PGSI–proportion-bet relationship, indicating differential effect across risk levels rather than uniform protection. Gender moderated MSG-POP effectiveness: specific framing types produced adverse patterns in female participants (Armstrong et al., 2018). STRUCT-STAKE showed minimal impact on non-problem gamblers but substantial expenditure effects for moderate- and problem-risk gamblers (Jackson et al., 2016). The LIMIT-LOSS hard-vs-soft-lock comparison documented an effectiveness-acceptability trade-off: hard locks produced greater gambling-activity reductions but higher short-term distress and lower satisfaction (Wohl et al., 2024).

Framework response. The Framework records target-population and friction-level fields for every tool, surfacing the moderation evidence the taxonomy made visible. The LIMIT-LOSS hard-vs-soft-lock distinction enters as separate entries: T013/T014 voluntary user-set (low friction), T015 mandatory regulatory ceiling (passive friction), T016 progress notification (low friction). The Wohl et al. effectiveness-acceptability trade-off sits across these as a documented design choice rather than a single recommended specification. EDU-VIDEO entries (T107–T110) reflect the moderation findings: T107 and T108 retained as universal-target tools (effective for low and moderate-risk per Hollingshead), T109 and T110 catalogued for high-risk awareness with framing dependencies. MSG-POP entries with documented risk-level reversals carry contraindication notes; tier-without-behavioural-targeting deployment is flagged as not supported in tool-selection logic.

Primary domains. MSG-POP, MSG-STATIC, EDU-VIDEO, STRUCT-STAKE, LIMIT-LOSS.

Pattern 7. Self-regulatory strategies as design resources

Spending and time limits are not interchangeable: the first RCT of assigned self-regulatory strategies in a real EGM environment found spending-limit assignment reduced harm while time-limit assignment increased it.

Spending vs time: not interchangeable. SELF-REG is the most recently emerging domain in the taxonomy (25 of 30 records post-2015) and the most uniformly Conceptual (25 of 30; 5 Validated, 0 Deployed). The first RCT to test assigned self-regulatory strategies in a real EGM environment produced a finding with direct implications for product design: assigned spending-limit strategies significantly reduced expenditure at follow-up, with strategy-use frequency mediating outcomes; assigned time-limit strategies were associated with increased spending and harms (Coelho et al., 2024). The contrast within a single study is notable: not all self-regulatory strategies are equivalent, and the common practice of presenting spending and time limits as interchangeable protective options is not supported by the evidence. The bust rate (exceeding self-set limit) was 45% overall and 83% among moderate- and high-risk gamblers (Rodda et al., 2019).

Spontaneous strategies as content templates. Across observational and qualitative records, pre-session spending limit-setting emerges as the most prevalent spontaneous strategy, rated helpful by 78.8% of users (Currie et al., 2020; Wood & Griffiths, 2015). Planning strategies received the highest helpfulness ratings overall, with no significant differences by severity, gender, or age (Rodda et al., 2018). Advance money limit-setting and activity balancing are associated with lower-risk gambling status (Hing et al., 2017). Time limit-setting received the lowest helpfulness rating among limit strategies and was associated with increased spending in longitudinal analysis (Currie et al., 2020; Coelho et al., 2024), consistent with the Coelho RCT finding. These player-generated strategies are largely absent from current platform features; they constitute design inputs supplied by naturalistic behavioural evidence that the taxonomy makes available.

Framework response. The Framework translates the spontaneous-strategy literature into player-facing tools by treating SELF-REG entries as content templates rather than standalone interventions. T137–T146 catalogue the strategies rated as most helpful in observational data: pre-session cash management (T137), advance budget commitment (T138), alternative-activity planning (T140), mood-state awareness (T141), and action-and-coping planning (T146). These supply the content base for low- and moderate-risk tool variants. T188 (nominated-person co-setting) extrapolates the third-party petition principle to voluntary limit management. The time-limit evidence shaped the limit-setting tools in the Framework directly: LIMIT-DEP defaults are spending-focused; LIMIT-TIME entries (T024–T029) are retained but flagged with the lower helpfulness ratings and the longitudinal increased-spending finding. Bust-rate evidence informs the limit-adherence support entries (T016 progress notification, T177 failed deposit counter).

Primary domains. SELF-REG, LIMIT-DEP, FEED-TRACK, GUIDE-LRGL, MSG-POP.

Pattern 8. Negative findings as design signals

Negative findings identify content formats and delivery conditions that are contraindicated rather than absent of evidence; the LIMIT-LOSS hard-vs-soft-lock RCT documents an effectiveness-acceptability trade-off that disappears in pooled summaries.

Twenty-one non-positive findings. Seventeen records are coded negative direction and four null, totalling 21 non-positive findings across the 410-record taxonomy (5.1%). Concentration patterns identify specific contraindications. MSG-STATIC has 4 negative records: responsibility-attribution framing activated reactance and increased rather than decreased gambling (Marko et al., 2022); women described static venue messages as ineffective and patronising (McCarthy et al., 2021); a poker website audit found static messages frequently non-functional or incorrectly placed (Khazaal et al., 2013). EDU-INFO has 3 negative and 1 null: neuroscience- and industry-framed content showed negligible or adverse effects on gambling attitudes (Myles et al., 2024); institutional messaging was rejected as irrelevant by players (Morvannou et al., 2020). EXCL-LONG has 1 negative and 1 null: no PGSI differences between those who renewed exclusion and those who discontinued (Pickering & Blaszczynski, 2022; Hing et al., 2022). SELF-REG has 1 negative: in-play

self-regulation attempts were associated with increased short-term gambling escalation, consistent with ego-depletion and reactance mechanisms (Coelho et al., 2024). LIMIT-LOSS contributes the effectiveness-acceptability trade-off: hard locks produced greater gambling reductions but higher short-term distress and lower satisfaction than soft locks (Wohl et al., 2024).

Framework response. The Framework treats negative findings as design constraints visible on each affected tool rather than reasons to omit. MSG-STATIC entries (T072–T075) are retained but reflect the negative evidence: T072 (“When the Fun Stops, Stop” generic slogan) is catalogued with the documented reactance and patronisation findings, while T074 (house-edge framing) and T075 (volatility statement) are flagged as the format variants supported by positive evidence. EDU-INFO entries (T103–T106) similarly retain the static-information format with caveats: T103 institutional facts page is documented as showing limited effects, with T104 explainer and T106 interactive cost calculator catalogued as the format variants with empirical support. SELF-REG entry T144 (self-monitoring journaling) carries the in-play self-regulation contraindication on its delivery specification, restricting the tool to between-session use. The LIMIT-LOSS hard-vs-soft-lock acceptability trade-off sits in the friction-level field (T013/T014 low-friction voluntary; T015 passive-friction mandatory).

Primary domains. MSG-STATIC, EDU-INFO, EXCL-LONG, SELF-REG, LIMIT-LOSS, LIMIT-DEP.